

Participantes:

Luis Gonzalez 2024-0509

Robert Yarel Zapata 2024-1020

Eimy Genao 2024-0543

Lyan Angomas 2024-191

PROYECTO LLM APLICACIONES

2025





TASK 1: Generacion de texto

Métrica: Perplexity

Dataset: medical-o1-reasoning-SFT / finepdfs

task	dataset	n_samples	model_id	size_hint	metric_name	metric_value
text_generatio	medical-o1 + f	500	openai-commi	small	perplexity	16.03433877
text_generatio	medical-o1 + f	500	gpt2-medium	medium	perplexity	12.3763627
text_generatio	medical-o1 + f	500	EleutherAI/gpt	small	perplexity	13.43912744
text_generatio	medical-o1 + f	500	EleutherAI/gpt	large	perplexity	8.616666063
inference_time_s	params_count	notes	Columna1	Columna2		
0.494907379	124439808	Perplexity manual, trunc=1024				
0.953192949	354823168	Perplexity manual, trunc=1024				
0.789925098	125198592	Perplexity manual, trunc=1024				
1.677183867	1315575808	Perplexity manual, trunc=1024				



TASK 2: Few Shot


Métrica: Accuracy

Dataset: SST-2

modelo	parametros	VRAM max (GB)	tiempo/ej (seg)	tokens/s	accuracy_4shot
FLAN-T5 Small	0.1B	2.3	0.1217	8.2	84.33%
FLAN-T5 Base	0.2B	2.66	0.1537	6.5	80.33%
FLAN-T5 Large	0.8B	3.71	0.3363	3	93.67%
BLOOM 560M	0.6B	3.62	0.2777	36	62.00%
GPT-Neo 1.3B	1.3B	4.93	0.4251	23.5	69.67%

f1_4shot	Validas	Dataset
86.30%	300	SST-2
83.75%	300	SST-2
94.04%	300	SST-2
71.78%	300	SST-2
61.28%	300	SST-2

Benchmark Few-Shot - Tiempo de Inferencia por Modelo




TASK 3: Zero Shot

Métrica: Accuracy

Dataset: SST-2

Modelo	Accuracy Zero-Shot (%)	Tiempo medio (s)	VRAM max (GB)	Validas	Dataset
FLAN-T5 Small (0.25B)	79.33	0.1609	3.8	300	SST-2
FLAN-T5 Base (0.78B)	89	0.1484	2.66	300	SST-2
FLAN-T5 Large (3B)	92.67	0.3764	3.71	300	SST-2
BLOOM 560M	48	0.2053	3.62	300	SST-2
GPT-Neo 1.3B	57.67	0.2863	4.93	300	SST-2




TASK 4: Traducción

Métrica: Bleu

Dataset: Helsinki-NLP/opus_books

task	direction	model_name	size_tag	num_parameters
translation	en->es	Helsinki-NLP/opus base		77943296
translation	en->es	Helsinki-NLP/opus big		234802176
translation	es->en	Helsinki-NLP/opus base		77943296
translation	es->en	vgaraujov/bart-base bart-base		139419648
approx_model_size_mb		bleu	num_samples	total_inference_time_sec
297.33		14.933	100	44.24
895.7		16.157	100	26
297.33		17.926	100	27.43
531.84		9.513	100	49.67
avg_time_per_sample_sec		status		
0.4424		ok		
0.26		ok		
0.2743		ok		
0.4967		ok		




TASK 5: Clasificación de texto

Métrica: Accuracy

Dataset: imdb

task	model	size	samples	total_time_s	avg_time_s	accuracy
sentiment-analysis	distilbert-base-uncased	small	300	324.0771332	1.080257111	0.88666667
sentiment-analysis	ProsusAI/finbert	small	300	686.8633811	2.289544604	1
sentiment-analysis	textattack/bert-base-uncased	medium	300	653.1351326	2.177117109	0.86333333
sentiment-analysis	roberta-base-uncased	medium	300	652.1554832	2.173851611	1
Dataset						
IMDB						
IMDB						
IMDB						
IMDB						



TASK 5: Resumen de texto

Métrica: Accuracy

Dataset: cnn_dailymail

model_id	nombre	n_samples	tiempo_total_s	tiempo_promedio_s	Dataset
facebook/bart	BART Base	5	157.9497409	31.58994818	cnn_dailymail
facebook/bart	BART Large CN	5	99.87061214	19.97412243	cnn_dailymail
t5-small	T5 Small	5	25.96866393	5.193732786	cnn_dailymail
t5-base	T5 Base	5	81.62084508	16.32416902	cnn_dailymail

ROUGE-L
0.1516711
0.29683714
0.26790627
0.23787812

